

# EDA(Algerian Forest Fires Dataset)

Rajesh Sekar

**Linkedin:** <https://lnkd.in/gu7QnXPV> (<https://lnkd.in/gu7QnXPV>).

**Github:** [https://github.com/Rajeshsekar1504/EDA\\_Algerian\\_Data\\_Set](https://github.com/Rajeshsekar1504/EDA_Algerian_Data_Set)  
([https://github.com/Rajeshsekar1504/EDA\\_Algerian\\_Data\\_Set](https://github.com/Rajeshsekar1504/EDA_Algerian_Data_Set)).

## EDA

- 1.Data Profiling
- 2.Statistical Analysis
- 3.Graphical Analysis

**Dataset:** <https://archive.ics.uci.edu/ml/datasets/Algerian+Forest+Fires+Dataset++>  
(<https://archive.ics.uci.edu/ml/datasets/Algerian+Forest+Fires+Dataset++>).

## Importing all the required the Libraries

In [1]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import plotly.express as px
import warnings

warnings.filterwarnings("ignore")
%matplotlib inline
pd.set_option('display.max_columns', 500)
```

## Importing dataset and cleaning data

In [2]:

```
df= pd.read_csv('Algerian_forest_fires_dataset_UPDATE.csv',header=1)
df.iloc[121:].head(4) #index 122,123 need to be removed from dataset
```

Out[2]:

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FW
121	30	09	2012	25	78	14	1.4	45	1.9	7.5	0.2	2.4	0.1
122	Sidi-Bel Abbes Region Dataset	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
123	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FW
124	01	06	2012	32	71	12	0.7	57.1	2.5	8.2	0.6	2.8	0.2

### Dropping rows which have no information

In [3]:

```
df.drop(index=[122,123], inplace=True) # dropping row 122,123 from dataset
df.reset_index(inplace=True)
df.drop('index', axis=1, inplace=True)
df.iloc[121:].head()
```

Out[3]:

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classe
121	30	09	2012	25	78	14	1.4	45	1.9	7.5	0.2	2.4	0.1	not fil
122	01	06	2012	32	71	12	0.7	57.1	2.5	8.2	0.6	2.8	0.2	not fil
123	02	06	2012	30	73	13	4	55.7	2.7	7.8	0.6	2.9	0.2	not fil
124	03	06	2012	29	80	14	2	48.7	2.2	7.6	0.3	2.6	0.1	not fil
125	04	06	2012	30	64	14	0	79.4	5.2	15.4	2.2	5.6	1	not fil

### Creating Region feature

In [4]:

```
### creating feature called Region 0 for Bejaia region and 1 for Sidi Bel-abbes region
df.loc[:122, 'Region']=0
df.loc[122:, 'Region']=1
df.iloc[120:].head(8)
```

Out[4]:

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classe
120	29	09	2012	26	80	16	1.8	47.4	2.9	7.7	0.3	3	0.1	not fire
121	30	09	2012	25	78	14	1.4	45	1.9	7.5	0.2	2.4	0.1	not fire
122	01	06	2012	32	71	12	0.7	57.1	2.5	8.2	0.6	2.8	0.2	not fire
123	02	06	2012	30	73	13	4	55.7	2.7	7.8	0.6	2.9	0.2	not fire
124	03	06	2012	29	80	14	2	48.7	2.2	7.6	0.3	2.6	0.1	not fire
125	04	06	2012	30	64	14	0	79.4	5.2	15.4	2.2	5.6	1	not fire
126	05	06	2012	32	60	14	0.2	77.1	6	17.6	1.8	6.5	0.9	not fire
127	06	06	2012	35	54	11	0.1	83.7	8.4	26.3	3.1	9.3	3.1	fire

## Datatypes and describe

In [5]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 244 entries, 0 to 243
Data columns (total 15 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   day             244 non-null   object  
 1   month          244 non-null   object  
 2   year           244 non-null   object  
 3   Temperature     244 non-null   object  
 4   RH              244 non-null   object  
 5   Ws              244 non-null   object  
 6   Rain            244 non-null   object  
 7   FFMC            244 non-null   object  
 8   DMC             244 non-null   object  
 9   DC              244 non-null   object  
10  ISI             244 non-null   object  
11  BUI             244 non-null   object  
12  FWI             244 non-null   object  
13  Classes         243 non-null   object  
14  Region          244 non-null   float64  
dtypes: float64(1), object(14)
memory usage: 28.7+ KB
```

## Describing the dataset

In [6]:

```
df.describe(include='all')
```

Out[6]:

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FW
count	244	244	244	244	244	244	244	244	244	244	244	244	244
unique	31	4	1	19	62	18	39	173	166	198	106	174	121
top	01	07	2012	35	64	14	0	88.9	7.9	8	1.1	3	0.4
freq	8	62	244	29	10	43	133	8	5	5	8	5	13
mean	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
std	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
min	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
25%	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
50%	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
75%	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
max	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

## Data Cleaning

In [7]:

```
# here it is visible that some columns have spaces in the names like RH, Ws
df.columns
```

Out[7]:

```
Index(['day', 'month', 'year', 'Temperature', ' RH', ' Ws', 'Rain ', 'FFMC',
      'DMC', 'DC', 'ISI', 'BUI', 'FWI', 'Classes ', 'Region'],
      dtype='object')
```

## stripping spaces from column names

In [8]:

```
df.columns = [col_name.strip() for col_name in df.columns]
df.columns
```

Out[8]:

```
Index(['day', 'month', 'year', 'Temperature', 'RH', 'Ws', 'Rain', 'FFMC',
      'DMC', 'DC', 'ISI', 'BUI', 'FWI', 'Classes', 'Region'],
      dtype='object')
```

In [9]:

```
### converting all feature values to string so that we can do data cleaning as shown below.  
df=df.astype(str)
```

In [10]:

```
### Some values in columns also have space  
for feature in ['Rain', 'FFMC', 'DMC', 'DC', 'ISI', 'BUI', 'FWI', 'Classes']:  
    df[feature] = df[feature].str.replace(" ", "")
```

In [11]:

```
## index no 165 for feature name FWI has value fire  
df[df['FWI']== 'fire'].index
```

Out[11]:

```
Int64Index([165], dtype='int64')
```

In [12]:

```
### replacing fire value witha float value  
df.loc[165, 'FWI'] = '0.1'
```

In [13]:

```
### replacing nan value wit fire to make data equal to the info given in dataset  
df[df['Classes']== 'nan'].index  
df.loc[165, 'Classes']='fire'
```

In [14]:

```
## encoding classes feature  
df['Classes']=df['Classes'].str.replace('notfire', '0')  
df['Classes']=df['Classes'].str.replace('fire', '1')
```

## Changing datatypes

In [15]:

```
### changing datatypes of features to numerical for numerical features as all are in object
datatype_convert={'day':'int64','month':'int64','year':'int64','Temperature':'int64','RH':'
'FFMC':'float64','DMC':'float64','DC':'float64','ISI':'float64','BUI':'float64','FWI':
'Classes':'int64','Region':'float64'}
df=df.astype(datatype_convert)
df.dtypes
```

Out[15]:

```
day          int64
month        int64
year         int64
Temperature  int64
RH           int64
Ws           int64
Rain         float64
FFMC         float64
DMC          float64
DC           float64
ISI          float64
BUI          float64
FWI          float64
Classes      int64
Region       float64
dtype: object
```

## Info about dataset and its attributes

1. The dataset includes 244 instances that regroup a data of two regions of Algeria,namely the Bejaia region located in the northeast of Algeria and the Sidi Bel-abbes region located in the northwest of Algeria.
2. 122 instances for each region.
3. The period from June 2012 to September 2012.
4. The dataset includes 11 attribues and 1 output attribue (class)
5. The 244 instances have been classified into fire (138 classes) and notfire (106 classes) classes.

## Attributes

1. Date : (DD/MM/YYYY) Day, month ('june' to 'september'), year (2012)

## Weather data observations

1. Temp : temperature noon (temperature max) in Celsius degrees: 22 to 42
2. RH : Relative Humidity in %: 21 to 90
3. Ws :Wind speed in km/h: 6 to 29
4. Rain: total day in mm: 0 to 16.8

## FWI Components

1. Fine Fuel Moisture Code (FFMC) index from the FWI system: 28.6 to 92.5
2. Duff Moisture Code (DMC) index from the FWI system: 1.1 to 65.9
3. Drought Code (DC) index from the FWI system: 7 to 220.4
4. Initial Spread Index (ISI) index from the FWI system: 0 to 18.5
5. Buildup Index (BUI) index from the FWI system: 1.1 to 68
6. Fire Weather Index (FWI) Index: 0 to 31.1
7. Classes: two classes, namely fire and not fire

In [16]:

```
df.shape
```

Out[16]:

```
(244, 15)
```

## Checking Null values

In [17]:

```
### checking for null values
```

```
df.isnull().sum()
```

Out[17]:

```
day          0
month        0
year         0
Temperature  0
RH           0
Ws           0
Rain         0
FFMC         0
DMC          0
DC           0
ISI          0
BUI          0
FWI          0
Classes      0
Region       0
dtype: int64
```

## Observation

1. There is no null value in dataset.
2. Total 244 rows and 15 columns is present.

## Numerical and continuous features

In [18]:

```
# categorical features

categorical_feature=[feature for feature in df.columns if df[feature].dtypes=='O']


#getting to know different categories in cateogrical features with its count.
for feature in categorical_feature:
    print(df.groupby(feature)['Region'].value_counts())
```

In [19]:

```
sns.countplot(data=df, x='Classes', hue='Region')
```

Out[19]:

<AxesSubplot:xlabel='Classes', ylabel='count'>



## Observation

1.It is evident that Sidi Bel-abbes region has more occurrence of fire than Bejaia region.

## Numerical features

In [20]:

```
### Getting list of numerical features
numerical_features=[feature for feature in df.columns if df[feature].dtypes!='O']
print(numerical_features)
```

```
['day', 'month', 'year', 'Temperature', 'RH', 'Ws', 'Rain', 'FFMC', 'DMC',  
'DC', 'ISI', 'BUI', 'FWI', 'Classes', 'Region']
```



In [21]:

```
### Getting unqiues values in each numerical features  
df[numerical_features].nunique()
```

Out[21]:

```
day          31  
month        4  
year         1  
Temperature  19  
RH           62  
Ws           18  
Rain        39  
FFMC        173  
DMC         166  
DC          198  
ISI         106  
BUI         174  
FWI         125  
Classes      2  
Region       2  
dtype: int64
```

## Seggregating discrete and continuous variables

### Discrete Numerical Features

In [22]:

```
#here the assumption to consider a feature discrete is that it should have less than 35 uni  
#considered continuous feature  
  
discrete_features = [feature for feature in numerical_features if len(df[feature].unique())  
discrete_features
```

Out[22]:


```
['day', 'month', 'year', 'Temperature', 'Ws', 'Classes', 'Region']
```

### Discrete Numerical Feature vs Target Feature

In [23]:

```
### this is bivariate analysis between target feature classes and discrete numerical featur

for feature in discrete_features:
    sns.countplot(data=df, x=feature, hue='Classes')
    plt.show()
```



## Observations

1. From day vs Classes plot it is visible that on almost all days the occurrence of fire is there, and its count is more than or equal to the count of no fire cases.
2. From month vs Classes plot it is visible that July and August month have more cases of occurrence of fire as compared to other two months of June and September where occurrence of fire is less as compared to no fire.
3. The month of August has highest no of cases of occurrence of fire.
4. Overall cases of occurrence of fire is more than the cases of no occurrence of fire.
5. From temperature vs Classes plot it is visible that temperature between 30 to 37 degree celcius have most no of cases of occurrence of fire.
6. From windspeed vs Classes plot it is visible that for wind speed between 13 to 19 Km/hr range there is most no of occurrence of fire.
7. From Region vs Class plot it is visible that in Bejaia region, the no of cases of occurrence of fire is less compared to no fire.
8. In Sidi Bel-abbes region the no of cases of occurrence of fire is more compared to no fire. Also Overall no of cases of occurrence of fire is more in Sidi Bel-abbes region as compared to Bejaia region.

## Continuous Numerical Features

In [24]:


```
continuous_features=[feature for feature in numerical_features if feature not in discrete_f  
print(continuous_features)
```

```
['RH', 'Rain', 'FFMC', 'DMC', 'DC', 'ISI', 'BUI', 'FWI']
```

## Distribution of Continuous Numerical Features

In [25]:

```
for feature in continuous_features:  
    sns.histplot(data=df,x=feature,kde=True,bins=30,color='blue')  
    plt.show()
```



## Observations


1. Relative humidity is following gaussian distribution
2. Rain,DMC,DC,ISI,BUI,FWI are following right skewed distribution(Log-Normal distribution)
3. FFMC feature follows left skewed distribution

## Checking for outliers

In [26]:

```
### excluding 'day','month','year','Region'
```

```
for feature in [feature for feature in numerical_features if feature not in ['day','month',  
sns.boxplot(data=df,y=feature,color='y')  
plt.title(feature)  
plt.show();
```



## Observations

1.Relative Humidity RH feature doesnt have outliers. 2.Temperature and FFMC have outliers in low boundary side. 3.Wind Speed, Ws has outliers on both sides(Upper and lower boundary). 4.Rain,DMC,ISI,BUI and FWI have outliers in upper boundary side.

## Correlation between each Numerical features

In [27]:

```
data = round(df[[feature for feature in numerical_features if feature not in ['day', 'month']  
data
```

Out[27]:

	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
Temperature	1.00	-0.65	-0.28	-0.33	0.68	0.48	0.37	0.61	0.46	0.56	0.52
RH	-0.65	1.00	0.24	0.22	-0.65	-0.41	-0.22	-0.69	-0.35	-0.57	-0.44
Ws	-0.28	0.24	1.00	0.17	-0.16	-0.00	0.08	0.02	0.03	0.03	-0.07
Rain	-0.33	0.22	0.17	1.00	-0.54	-0.29	-0.30	-0.35	-0.30	-0.32	-0.38
FFMC	0.68	-0.65	-0.16	-0.54	1.00	0.60	0.50	0.74	0.59	0.69	0.77
DMC	0.48	-0.41	-0.00	-0.29	0.60	1.00	0.88	0.67	0.98	0.87	0.58
DC	0.37	-0.22	0.08	-0.30	0.50	0.88	1.00	0.50	0.94	0.74	0.51
ISI	0.61	-0.69	0.02	-0.35	0.74	0.67	0.50	1.00	0.64	0.91	0.74
BUI	0.46	-0.35	0.03	-0.30	0.59	0.98	0.94	0.64	1.00	0.86	0.58
FWI	0.56	-0.57	0.03	-0.32	0.69	0.87	0.74	0.91	0.86	1.00	0.71
Classes	0.52	-0.44	-0.07	-0.38	0.77	0.58	0.51	0.74	0.58	0.71	1.00


Heatmap to visualise the Correlation

In [28]:

```
### Plotting heatmap for visualising the correlation between features
sns.set(rc={'figure.figsize':(15,10)})
sns.heatmap(data=data,annot=True, vmin=-1, vmax=1)
```

Out[28]:

&lt;AxesSubplot:&gt;



## Note(For both positive and negative side)

1. Correlation coefficients between 0.9 and 1.0,very highly correlated.
2. Correlation coefficients between 0.7 and 0.9, highly correlated.
3. Correlation coefficients between 0.5 and 0.7, moderately correlated.
4. Correlation coefficients between 0.3 and 0.5, low correlation.
5. Correlation coefficients less than 0.3, little correlation.

## Observations

1. Very highly Correlated feature:DMC-BUI,DC-BUI,ISI-FWI
2. Highly correlated features:FFMC-ISI,DC-DMC,FWI-DMC,FWI-DC,FWI-BUI Note: Features with very high and high correlation are more linearly dependent and hence have almost the same effect on the dependent variable. So, we can drop one of the two features.

## Feature VS target

## day

In [29]:

```
sns.scatterplot(data=df,x='day',y='Temperature',hue='Classes')
```

Out[29]:

<AxesSubplot:xlabel='day', ylabel='Temperature'>



## Observation

1. Most cases of fire occur for temperature more than 30 degree celcius

In [30]:

```
df
```

Out[30]:

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Class
0	1	6	2012	29	57	18	0.0	65.7	3.4	7.6	1.3	3.4	0.5	
1	2	6	2012	29	61	13	1.3	64.4	4.1	7.6	1.0	3.9	0.4	
2	3	6	2012	26	82	22	13.1	47.1	2.5	7.1	0.3	2.7	0.1	
3	4	6	2012	25	89	13	2.5	28.6	1.3	6.9	0.0	1.7	0.0	
4	5	6	2012	27	77	16	0.0	64.8	3.0	14.2	1.2	3.9	0.5	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	
239	26	9	2012	30	65	14	0.0	85.4	16.0	44.5	4.5	16.9	6.5	
240	27	9	2012	28	87	15	4.4	41.1	6.5	8.0	0.1	6.2	0.0	
241	28	9	2012	27	87	29	0.5	45.9	3.5	7.9	0.4	3.4	0.2	
242	29	9	2012	24	54	18	0.1	79.7	4.3	15.2	1.7	5.1	0.7	
243	30	9	2012	24	64	15	0.2	67.3	3.8	16.5	1.2	4.8	0.5	

244 rows × 15 columns



## Month




In [31]:

```
sns.scatterplot(data=df,x='month',y='Temperature',hue='Classes')
```

Out[31]:

<AxesSubplot:xlabel='month', ylabel='Temperature'>



## Observations

1. July and august have more cases of fire as compared to no fire.
2. june and september have more cases of no fire as compared to fire.


## Region

In [32]:

```
sns.scatterplot(data=df,x='Region',y='Temperature',hue='Classes')
```

Out[32]:

<AxesSubplot:xlabel='Region', ylabel='Temperature'>

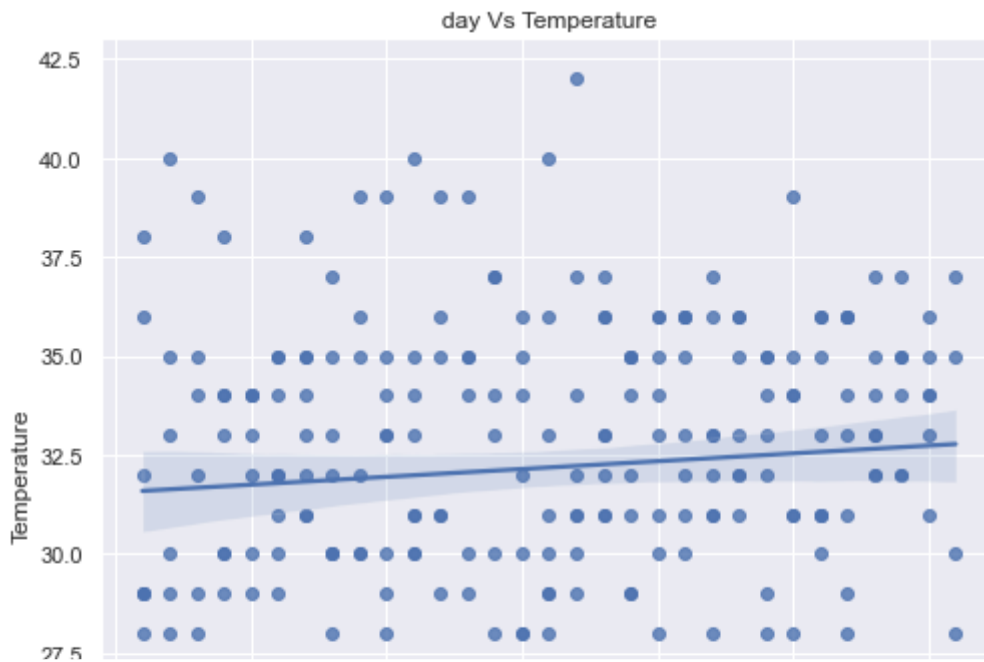


## Observations

1. In Bejaia region, the no of cases of occurance of fire is less compared to no of cases of occurance of no fire.
2. In Sidi Bel-abbes region the no of cases of occurance of fire is more compared to no fire.
3. Also Overall no of cases of occurance of fire is more in Sidi Bel-abbes region as compared to Bejaia region.

In [33]:

```
##### shaded region is basically with respect to ridge and lasso (lambda)
for feature in [feature for feature in df.columns if feature not in ['Temp']]:
    sns.set(rc={'figure.figsize':(8,8)})
    sns.regplot(x=df[feature], y=df['Temperature'])
    plt.xlabel(feature)
    plt.ylabel('Temperature')
    plt.title('{} Vs Temperature'.format(feature))
    plt.show();
```



## Final Report

1. very highly correlated features:DMC-BUI,ISI-FWI
2. Highly correlated features:FFMC-ISI,DC-DMC,FWI-DMC,FWI-DC,FWI-BUI
3. Temperature between 30 to 37 degree celcius have most no of cases of occurance of fire.
4. Wind speed between 13 to 19Km/hr range there is most no of occurance of fire.
5. Almost all cases of occurance of fires is for region having rain less than 1mm i.e, dry regions are more prone to forest fires.
6. For FFMC(Fine Fuel Moisture Code) greater that 80, almost all cases of fire is reported.
7. DMC (Duff Moisture Code) >30 and DC (Drought code) >100, almost all cases of occurance of fire reported, this means drought affected areas are more prone to forrest fires.
8. In Bejaia region, the no of cases of occurance of fire is less compared to no of cases of occurance of no fire.
9. In Sidi Bel-abbes region the no of cases of occurance of fire is more compared to no fire.
10. Also Overall no of cases of occurance of fire is more in Sidi Bel-abbes region as compared to Bejaia region.
11. Most no of cases of fire occured are in the month of august and least no of cases of fire occured is in month of september.
12. July and august have more cases of fire as compared to no fire.
13. June and september have more cases of no fire as compared to fire.
14. Relative Humidity, RH feature doesnt have outliers whereas Temperature, FFMC, wind speed, Rain, DMC,DC, ISI, BUI and FWI have outilers.
15. There is no null vales in dataset.

**Note** EDA and basic feature engineering is done its time to seperate independent and dependent features.

1. For demonstrating linear regression taking Temperature as Dependent feature.
2. dropping year feature as dataset contains only 2012 year

In [34]:

```
df.drop('year',axis=1,inplace=True)
```

In [35]:

```
df['Temp']=df['Temperature']
```

In [36]:

```
df.drop('Temperature',axis=1,inplace=True)
```

In [37]:

```
df
```

Out[37]:

	day	month	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes	Region	Temp
0	1	6	57	18	0.0	65.7	3.4	7.6	1.3	3.4	0.5	0	0.0	29
1	2	6	61	13	1.3	64.4	4.1	7.6	1.0	3.9	0.4	0	0.0	29
2	3	6	82	22	13.1	47.1	2.5	7.1	0.3	2.7	0.1	0	0.0	26
3	4	6	89	13	2.5	28.6	1.3	6.9	0.0	1.7	0.0	0	0.0	25
4	5	6	77	16	0.0	64.8	3.0	14.2	1.2	3.9	0.5	0	0.0	27
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
239	26	9	65	14	0.0	85.4	16.0	44.5	4.5	16.9	6.5	1	1.0	30
240	27	9	87	15	4.4	41.1	6.5	8.0	0.1	6.2	0.0	0	1.0	28
241	28	9	87	29	0.5	45.9	3.5	7.9	0.4	3.4	0.2	0	1.0	27
242	29	9	54	18	0.1	79.7	4.3	15.2	1.7	5.1	0.7	0	1.0	24
243	30	9	64	15	0.2	67.3	3.8	16.5	1.2	4.8	0.5	0	1.0	24

244 rows × 14 columns

## Starting Model Building Preparatin

In [38]:

```
### Getting Independent features in dataset and Dependent feature in Series object
```

In [39]:

```
df.head()
```

Out[39]:

	day	month	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes	Region	Temp
0	1	6	57	18	0.0	65.7	3.4	7.6	1.3	3.4	0.5	0	0.0	29
1	2	6	61	13	1.3	64.4	4.1	7.6	1.0	3.9	0.4	0	0.0	29
2	3	6	82	22	13.1	47.1	2.5	7.1	0.3	2.7	0.1	0	0.0	26
3	4	6	89	13	2.5	28.6	1.3	6.9	0.0	1.7	0.0	0	0.0	25
4	5	6	77	16	0.0	64.8	3.0	14.2	1.2	3.9	0.5	0	0.0	27

In [40]:

```
### X independent features and y dependent feature
X = df.iloc[:, :-1]
y = df.iloc[:, -1]
```

In [41]:

```
X.head()
```

Out[41]:

	day	month	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes	Region
0	1	6	57	18	0.0	65.7	3.4	7.6	1.3	3.4	0.5	0	0.0
1	2	6	61	13	1.3	64.4	4.1	7.6	1.0	3.9	0.4	0	0.0
2	3	6	82	22	13.1	47.1	2.5	7.1	0.3	2.7	0.1	0	0.0
3	4	6	89	13	2.5	28.6	1.3	6.9	0.0	1.7	0.0	0	0.0
4	5	6	77	16	0.0	64.8	3.0	14.2	1.2	3.9	0.5	0	0.0

In [42]:

```
y.head()
```

Out[42]:

```
0    29
1    29
2    26
3    25
4    27
Name: Temp, dtype: int64
```

## Splitting data into Training and Test data

In [43]:

```
## Splitting data into Training and Test da
from sklearn.model_selection import train_test_split
```

In [44]:

```
### random state train test split will be same with all people using random state =42
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.33,random_state=42)
```

In [45]:

```
X_train.head()
```

Out[45]:

	day	month	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes	Region
<b>114</b>	23	9	54	11	0.5	73.7	7.9	30.4	1.2	9.6	0.7	0	0.0
<b>65</b>	5	8	65	13	0.0	86.8	11.1	29.7	5.2	11.5	6.1	1	0.0
<b>132</b>	11	6	42	21	0.0	90.6	18.2	30.5	13.4	18.0	16.7	1	1.0
<b>207</b>	25	8	40	18	0.0	92.1	56.3	157.5	14.3	59.5	31.1	1	1.0
<b>162</b>	11	7	56	15	2.9	74.8	7.1	9.5	1.6	6.8	0.8	0	1.0

In [46]:

```
y_train.head()
```

Out[46]:

```
114    32
65     34
132    31
207    34
162    34
Name: Temp, dtype: int64
```

In [47]:

```
X_test.head()
```

Out[47]:

	day	month	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes	Region
<b>24</b>	25	6	64	15	0.0	86.7	14.2	63.8	5.7	18.3	8.4	1	0.0
<b>6</b>	7	6	54	13	0.0	88.2	9.9	30.5	6.4	10.9	7.2	1	0.0
<b>153</b>	2	7	48	16	0.0	87.6	7.9	17.8	6.8	7.8	6.4	1	1.0
<b>211</b>	29	8	53	17	0.5	80.2	20.7	149.2	2.7	30.6	5.9	1	1.0
<b>198</b>	16	8	41	10	0.1	92.0	22.6	65.1	9.5	24.2	14.8	1	1.0

In [48]:

```
y_test.head()
```

Out[48]:

```
24      31
6       33
153     33
211     35
198     40
Name: Temp, dtype: int64
```

In [49]:

```
### both will have same shape
X_train.shape,y_train.shape
```

Out[49]:

```
((163, 13), (163,))
```

In [50]:

```
### both will have same shape
X_test.shape,y_test.shape
```

Out[50]:

```
((81, 13), (81,))
```

## Feature Engineering

### Standardisation/feature scaling the dataset

In [51]:

```
from sklearn.preprocessing import StandardScaler
```

In [52]:

```
### creating a StandardScaler object
scaler = StandardScaler()
scaler
```

Out[52]:

```
StandardScaler()
```

In [53]:

```
### Using fit_transform to standardise train data
X_train = scaler.fit_transform(X_train)
X_train
```

Out[53]:

```
array([[ 0.84447703,  1.3826723 , -0.60257784, ..., -0.8196431 ,
        -1.04390785, -0.99388373],
       [-1.19310159,  0.48116996,  0.14460201, ..., -0.08219052,
        0.95793896, -0.99388373],
       [-0.51390872, -1.32183472, -1.41768313, ...,  1.36540157,
        0.95793896,  1.0061539 ],
       ...,
       [-1.64589683,  1.3826723 ,  0.89178186, ..., -0.90158227,
        -1.04390785, -0.99388373],
       [ 1.41047108, -0.42033238, -0.39880152, ...,  0.31384882,
        0.95793896,  1.0061539 ],
       [-0.51390872,  1.3826723 ,  0.9597073 , ..., -0.87426921,
        -1.04390785, -0.99388373]])
```

In [54]:

```
### Here using only transform to avoid data leakage
### (training mean and training std will be used for standardisation of test when we use tr
X_test=scaler.transform(X_test)
X_test
```

Out[54]:

```
array([[ 1.07087465, -1.32183472,  0.07667657, ...,  0.23190965,
        0.95793896, -0.99388373],
       [-0.96670396, -1.32183472, -0.60257784, ...,  0.0680313 ,
        0.95793896, -0.99388373],
       [-1.53269802, -0.42033238, -1.01013048, ..., -0.04122093,
        0.95793896,  1.0061539 ],
       ...,
       [ 1.29727227, -0.42033238, -1.01013048, ...,  1.17421016,
        0.95793896, -0.99388373],
       [-1.3063004 , -1.32183472,  0.07667657, ..., -0.77867351,
        -1.04390785,  1.0061539 ],
       [ 1.29727227, -1.32183472, -0.5346524 , ...,  0.7235447 ,
        0.95793896,  1.0061539 ]])
```

## Model Buliding

## Linear Regression

In [55]:

```
from sklearn.linear_model import LinearRegression
```



In [56]:

```
### creating Linear regression model
linear_reg = LinearRegression()
linear_reg
```

Out[56]:

LinearRegression()

In [57]:

```
### passing training data(X and y ) to the model
linear_reg.fit(X_train,y_train)
```

Out[57]:

LinearRegression()

In [58]:

```
## Printing co-efficients and intercept of best fit hyperplane
print("1. Co-efficients of independent features is {}".format(linear_reg.coef_))
print("2. Intercept of best fit hyperplane is {}".format(linear_reg.intercept_))
```

```
1. Co-efficients of independent features is [-0.62994684 -0.33080692 -0.9681
523 -0.55769053  0.23645285  1.90585623
 0.93380592  1.17296981  0.044581   -1.35995788  0.19772494 -0.25230922
 0.08345626]
2. Intercept of best fit hyperplane is 31.98159509202454
```

## Using model to get predictions of test data

In [59]:

```
linear_reg_pred = linear_reg.predict(X_test)
linear_reg_pred
```

Out[59]:

```
array([32.86982262, 34.97907511, 34.71895423, 32.93220734, 36.64866482,
       32.00281859, 35.27819508, 28.49312857, 31.84450923, 29.27704091,
       29.06704133, 33.07364481, 32.4667427 , 32.7008168 , 34.32599535,
       31.80453584, 37.01042617, 25.23211237, 32.73196597, 33.38253854,
       31.55571716, 28.30699286, 34.23615097, 29.30603632, 36.93126913,
       24.98756128, 33.51228222, 33.57587507, 33.35705604, 35.40329932,
       33.767112 , 31.85221582, 32.40507656, 33.11736397, 32.44972087,
       31.46599605, 30.34784931, 34.2239929 , 32.37589956, 21.74277219,
       33.82900884, 34.85103093, 31.20651563, 24.69868309, 36.17424894,
       32.81796744, 31.22635993, 30.67357508, 35.1950892 , 34.29311524,
       36.98975313, 30.97884914, 30.95678802, 34.6655222 , 33.46814569,
       32.38222097, 36.65227179, 30.589826 , 30.97603618, 36.10290928,
       33.94615809, 28.43783118, 33.17776773, 31.78923636, 31.99593987,
       24.12810241, 33.39123143, 29.76320324, 36.80847578, 34.30376941,
       33.61696277, 31.49444654, 33.44085947, 34.43788629, 35.59708798,
       31.17211416, 32.72579793, 32.96039667, 35.20161022, 33.43024933,
       33.69316482])
```

## Validating model using assumptions of Linear regression

### Linear relationship


1. Test truth data and Predicted data should follow linear relationship.
2. This is an indication of a good model.

In [60]:

```
plt.scatter(x = y_test, y = linear_reg_pred)
plt.xlabel("Test truth data")
plt.ylabel("Predicted data")
```

Out[60]:

Text(0, 0.5, 'Predicted data')



### Residual distribution

1. Residuals should follow normal distribution.
2. if residuals follow normal distribution, it indicates we have a good model.

In [61]:

```
residual_linear_reg=y_test-linear_reg_pred  
residual_linear_reg.head()
```

Out[61]:


```
24    -1.869823  
6     -1.979075  
153   -1.718954  
211    2.067793  
198    3.351335  
Name: Temp, dtype: float64
```

In [62]:

```
sns.displot(x=residual_linear_reg, kind='kde')
```

Out[62]:

<seaborn.axisgrid.FacetGrid at 0x126c8ff1820>



## Uniform distribution


1. Residuals vs Predictions should follow a uniform distribution.
2. If Residuals vs Predictions follow uniform distribution, it indicates we have a good model.

In [63]:

```
plt.scatter(x=linear_reg_pred,y=residual_linear_reg)
plt.xlabel('Predictions')
plt.ylabel('Residuals')
```

Out[63]:

Text(0, 0.5, 'Residuals')



## Performance Matrix

## Cost function values

In [64]:

```
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
```

## ## MSE, MAE and RMSE

In [65]:

```
print("Mean squared error is {}".format(round(mean_squared_error(y_test,linear_reg_pred),2))
print("Mean absolute error is {}".format(round(mean_absolute_error(y_test,linear_reg_pred),
print("Root mean squared error is {}".format(round(np.sqrt(mean_squared_error(y_test,linear
```

Mean squared error is 5.25  
Mean absolute error is 1.81  
Root mean squared error is 2.29

## ### R Square and Adjusted R Square values

In [66]:

```
from sklearn.metrics import r2_score
```

In [67]:

```
linear_reg_r2_score=r2_score(y_test,linear_reg_pred)
print("Our Linear regression model has {} % accuracy".format(round(linear_reg_r2_score*100,

linear_reg_r2_score=1-((1-linear_reg_r2_score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]
print("Adjusted R square accuracy is {} percent".format(round(linear_reg_r2_score*100,2)))
```

Our Linear regression model has 51.089 % accuracy  
Adjusted R square accuracy is 41.6 percent

## Ridge Regression

In [68]:

```
from sklearn.linear_model import Ridge
```

In [69]:

```
## Creating Ridge regression model
ridge_reg = Ridge()
ridge_reg
```

Out[69]:

Ridge()

In [70]:

```
### Passing training data(X and y) to the model
ridge_reg.fit(X_train,y_train)
```

Out[70]:

Ridge()

In [71]:

```
### Printing co-efficients and intercept of best fit hyperplane
print("1. Co-efficients of independent features is {}".format(ridge_reg.coef_))
print("2. Intercept of best fit hyperplane is {}".format(ridge_reg.intercept_))
```

```
1. Co-efficients of independent features is [-0.61752995 -0.3207458 -0.9821
8457 -0.55467826  0.21315492  1.84131702
 0.40094067  0.87094221  0.11396182 -0.48465778  0.08466793 -0.24063909
 0.09187935]
2. Intercept of best fit hyperplane is 31.98159509202454
```

## Using model to get predictions of test data

In [72]:

```
ridge_reg_pred = ridge_reg.predict(X_test)
ridge_reg_pred
```

Out[72]:

```
array([32.85982748, 34.9149207 , 34.6801255 , 32.92998132, 36.61056862,
       32.05917754, 35.25499575, 28.51988807, 31.83679288, 29.28276684,
       29.06439442, 33.15037313, 32.44677748, 32.7322483 , 34.35814975,
       31.78754571, 36.91567124, 25.31324295, 32.67492302, 33.35526777,
       31.49765658, 28.29644553, 34.22316335, 29.30830796, 36.93224783,
       25.06526261, 33.48461631, 33.57641555, 33.35387649, 35.32860357,
       33.76792408, 31.83362797, 32.37993338, 33.16010038, 32.42109455,
       31.5078178 , 30.28919718, 34.25440049, 32.3007279 , 21.86171205,
       33.83192673, 34.81983629, 31.24261825, 24.76831106, 36.10419592,
       32.77320818, 31.21807666, 30.69260083, 35.1742616 , 34.29037291,
       36.93483074, 30.9447582 , 30.97755205, 34.71719979, 33.43904851,
       32.52070875, 36.64973602, 30.62216011, 30.9696123 , 36.10416977,
       33.88939183, 28.47020463, 33.13493675, 31.7769902 , 32.00695307,
       24.17731957, 33.37004249, 29.76297127, 36.78001537, 34.44394437,
       33.58457247, 31.47437138, 33.41857278, 34.45102238, 35.62964268,
       31.16568509, 32.68759611, 32.91480612, 35.20066129, 33.38322692,
       33.67748654])
```

## Validating Model using assumptions of Ridge regression

### Linear relationship


1. Test truth data and Predicted data should follow linear relationship.
2. This is an indication of a good model.

In [73]:

```
plt.scatter(x=y_test,y=ridge_reg_pred)
plt.xlabel("Test truth data")
plt.ylabel("Predicted data")
```

Out[73]:

Text(0, 0.5, 'Predicted data')



## Residual distribution

1. Residuals should follow normal distribution.
2. If residuals follow normal distribution, it indicates we have a good model.

In [74]:

```
residual_ridge_reg = y_test-ridge_reg_pred  
residual_ridge_reg.head()
```

Out[74]:


```
24    -1.859827  
6     -1.914921  
153   -1.680125  
211    2.070019  
198    3.389431  
Name: Temp, dtype: float64
```

In [75]:

```
sns.displot(x=residual_ridge_reg, kind = 'kde')
```

Out[75]:

<seaborn.axisgrid.FacetGrid at 0x126cb105d60>



## Uniform distribution

1. Residuals vs Predictions should follow a uniform distribution.
2. If Residuals vs Predictions follow uniform distribution , it indicates we have a good model.




In [76]:

```
plt.scatter(x=ridge_reg_pred,y=residual_ridge_reg)
plt.xlabel('Predictions')
plt.ylabel('Residuals')
```

Out[76]:

Text(0, 0.5, 'Residuals')



## Performance Matrix

Cost function values

**MSE, MAE, RMSE**

In [77]:

```
print("Mean squared error is {}".format(round(mean_squared_error(y_test,ridge_reg_pred),2))
print("Mean absolute error is {}".format(round(mean_absolute_error(y_test,ridge_reg_pred),2)
print("Root Mean squared error is {}".format(round(np.sqrt(mean_squared_error(y_test,ridge_
```

Mean squared error is 5.19  
Mean absolute error is 1.8  
Root Mean squared error is 2.28

## R Square and Adjusted R Square values

In [82]:

```
ridge_reg_r2_score = r2_score(y_test, ridge_reg_pred)
print("Our Ridge regression model has {} % accuracy".format(round(ridge_reg_r2_score*100,3)
ridge_reg_adj_r2_score = 1-((1-ridge_reg_r2_score)*(len(y_test)-1)/(len(y_test)-X_test.shap
print("Adjusted R square accuracy is {} percent".format(round(ridge_reg_adj_r2_score*100,2))
```

Our Ridge regression model has 51.709 % accuracy  
Adjusted R square accuracy is 42.34 percent

## Lasso Regression

In [83]:

```
from sklearn.linear_model import Lasso
```

In [84]:

```
## Creating Lasso regression model
lasso_reg = Lasso()
lasso_reg
```

Out[84]:

Lasso()

In [85]:

```
### Passing training data(X and y) to the model
lasso_reg.fit(X_train,y_train)
```

Out[85]:

Lasso()

In [86]:

```
### Printing co-efficients and intercept of best fit hyperplane
print("1. Co-efficients of independent feautures is {}".format(lasso_reg.coef_))
print("2. Intercept of best fit hyperplane is {}".format(lasso_reg.intercept_))
```

```
1. Co-efficients of independent feautures is [-0.          -0.          -0.623
24302 -0.          -0.          1.25581509
  0.          0.          0.          0.          0.          0.
  0.          ]
2. Intercept of best fit hyperplane is 31.98159509202454
```

## Using model to get predictions of test data

In [87]:

```
lasso_reg_pred = lasso_reg.predict(X_test)
lasso_reg_pred
```

Out[87]:

```
array([32.78381104, 33.3358205 , 33.53835729, 32.69192045, 34.21212444,
       31.67725854, 34.06518855, 28.84685412, 30.99078013, 30.10392027,
       31.06631475, 32.42020469, 32.80398907, 32.31726957, 33.37068778,
       32.46976122, 34.57875298, 27.38502889, 32.29240264, 33.12192792,
       31.62499111, 29.79900395, 33.65042591, 30.31559056, 34.62053146,
       28.73534108, 32.5527626 , 32.85187888, 32.99739235, 33.91411932,
       33.25917474, 31.27076256, 32.79238897, 32.82670067, 32.58151856,
       32.71432078, 31.4254765 , 33.30564213, 31.04995877, 28.31033379,
       32.79596672, 32.94981382, 32.61305239, 27.48796401, 34.24588058,
       32.5549849 , 31.88399562, 30.89031162, 34.19107957, 33.50571231,
       34.39861652, 31.81203876, 31.79074958, 32.97252542, 33.46393383,
       32.11393291, 34.88311373, 32.60669677, 31.06489233, 34.43070593,
       33.14932842, 31.16313854, 33.31119788, 32.23069043, 32.12966634,
       27.92377153, 33.05552681, 31.68837004, 34.34770454, 34.11554496,
       32.95092497, 31.93212973, 33.10008317, 33.54804636, 34.5256187 ,
       31.88399562, 33.11557229, 32.30511389, 33.7699613 , 32.15762242,
       33.37068778])
```

## Validating model using assumptions of Lasso regression

### Linear relationship


1. Test truth data and predicted data should follow linear relationship.
2. This is an indication of a good model.

In [89]:

```
plt.scatter(x=y_test, y=lasso_reg_pred)
plt.xlabel("Test truth data")
plt.ylabel("Predicted data")
```

Out[89]:

Text(0, 0.5, 'Predicted data')



## Residual distribution

1. Residuals should follow normal distribution
2. If residuals follow normal distribution, it indicates we have a good model.

In [91]:

```
residual_lasso_reg=y_test-lasso_reg_pred  
residual_lasso_reg.head()
```

Out[91]:


```
24    -1.783811  
6      -0.335821  
153   -0.538357  
211    2.308080  
198    5.787876  
Name: Temp, dtype: float64
```

In [92]:

```
sns.displot(x=residual_lasso_reg, kind='kde')
```

Out[92]:

<seaborn.axisgrid.FacetGrid at 0x126cc6d27f0>



## Uniform distribution


1. Residuals vs Predictions should follow a uniform distribution.
2. If Residuals vs Predictions follow uniform distribution, it indicates we have a good model.

In [93]:

```
plt.scatter(x=lasso_reg_pred, y=residual_lasso_reg)
plt.xlabel('Predictions')
plt.ylabel('Residuals')
```

Out[93]:

Text(0, 0.5, 'Residuals')



## Performance Matrix

## Cost function values

MSE, MAE and RMSE

In [96]:

```
print("Mean squared error is {}".format(round(mean_squared_error(y_test, lasso_reg_pred),2))
print("Mean absolute error is {}".format(round(mean_absolute_error(y_test, lasso_reg_pred),
print("Root Mean squared error is {}".format(round(np.sqrt(mean_squared_error(y_test, lasso
```

Mean squared error is 6.09

Mean absolute error is 2.0

Root Mean squared error is 2.47

## R Square and Adjusted R Square values

In [97]:

```
lasso_reg_r2_score=r2_score(y_test, lasso_reg_pred)
print("Our Lasso regression model has {} % accuracy".format(round(lasso_reg_r2_score*100,3))
lasso_reg_adj_r2_score=1-((1-lasso_reg_r2_score)*(len(y_test)-1)/(len(y_test)-X_test.shape[
print("Adjusted R square accuracy is {}".format(round(lasso_reg_r2_score*100,2)))
```

Our Lasso regression model has 43.342 % accuracy

Adjusted R square accuracy is 43.34

## Elastic-Net Regression

In [98]:

```
from sklearn.linear_model import ElasticNet
```

In [99]:

```
## creating Elastic-Net regression model
elastic_reg = ElasticNet()
elastic_reg
```

Out[99]:

ElasticNet()

In [100]:

```
### Passing training data(X and y) to the model
elastic_reg.fit(X_train, y_train)
```

Out[100]:

ElasticNet()

In [101]:

```
### Printing co-efficients and intercept of best fit hyperplane
print("1. Co-efficients of independent features is {}".format(elastic_reg.coef_))
print("2. Intercept of best fit hyperplane is {}".format(elastic_reg.intercept_))
```

```
1. Co-efficients of independent features is [-0.          -0.          -0.6880
8933 -0.10544712 -0.00834786  0.85162206
 0.10376148  0.          0.23158765  0.02547021  0.15362153  0.07372069
 0.          ]
2. Intercept of best fit hyperplane is 31.98159509202454
```

## Using model to get predictions of test data

In [102]:

```
elastic_reg_pred=elastic_reg.predict(X_test)
elastic_reg_pred
```

Out[102]:

```
array([32.70014869, 33.29910099, 33.41026626, 32.61092932, 34.7047485 ,
       31.58360838, 34.21527053, 29.02563256, 30.73347022, 30.11039166,
       30.41277398, 32.24275851, 32.49689882, 32.11572726, 33.82943086,
       32.26602144, 35.60101706, 27.89264401, 32.12951491, 32.95265792,
       31.02735367, 29.65251053, 33.71392821, 30.09882338, 35.36763797,
       28.87850096, 32.39310489, 32.86323328, 32.70140428, 34.17896096,
       33.34379291, 31.19918174, 32.72501691, 33.05000354, 32.11073799,
       32.39596102, 30.82655691, 33.29618332, 31.03905163, 27.97071052,
       32.81078878, 32.83085009, 32.18944481, 27.83881794, 34.71347157,
       32.40467835, 31.69547324, 30.73847724, 34.32027173, 33.61283179,
       35.5799204 , 31.38989764, 31.57088417, 33.58592157, 33.61999323,
       32.2560815 , 36.04767586, 32.00594323, 30.83622131, 34.89715148,
       33.0854042 , 30.78738109, 33.49083492, 31.9555529 , 31.71431021,
       27.98432489, 32.99835463, 31.35802614, 35.34742765, 34.15121761,
       32.70498873, 31.52572086, 33.29706361, 33.77686044, 35.16781422,
       31.6287495 , 33.16854289, 32.00367904, 34.13279424, 31.72123893,
       33.52323673])
```

## Validating model using assumptions of Elastic-Net regression

### Linear relationship

1. Test truth data and Predicted data should follow linear relationship.
2. this is an indication of a good model.




In [103]:

```
plt.scatter(x=y_test,y=elastic_reg_pred)
plt.xlabel("Test truth data")
plt.ylabel("Predicted data")
```

Out[103]:

Text(0, 0.5, 'Predicted data')



## Residual distribution

1. Residuals should follow normal distribution.
2. if residuals follow normal distribution, it indicates we have a good model.

In [105]:

```
residual_elastic_reg = y_test-elastic_reg_pred  
residual_elastic_reg.head()
```

Out[105]:


```
24    -1.700149  
6     -0.299101  
153   -0.410266  
211    2.389071  
198    5.295251  
Name: Temp, dtype: float64
```

In [106]:

```
sns.displot(x=residual_elastic_reg, kind='kde')
```

Out[106]:

<seaborn.axisgrid.FacetGrid at 0x126cd4cfc70>



## Uniform distribution


1. Residuals vs Predictions should follow a uniform distribution.
2. If Residuals vs Predictions follow uniform distribution, it indicates we have a good model.

In [107]:

```
plt.scatter(x=elastic_reg_pred, y = residual_elastic_reg)
plt.xlabel('Predictions')
plt.ylabel('Residuals')
```

Out[107]:

Text(0, 0.5, 'Residuals')



## Performance Matrix

### Cost function values

### MSE, MAE and RMSE

In [108]:

```
print("Mean squared error is {}".format(round(mean_squared_error(y_test, elastic_reg_pred),
print("Mean absolute error is {}".format(round(mean_absolute_error(y_test,elastic_reg_pred)
print("Root Mean squared error is {}".format(round(np.sqrt(mean_squared_error(y_test,elasti
```

Mean squared error is 5.39  
Mean absolute error is 1.85  
Root Mean squared error is 2.32

## R Square and Adjusted R Square values

In [109]:

```
elastic_reg_r2_score=r2_score(y_test, elastic_reg_pred)
print("Our Elatic-Net regression model has {} % accuracy".format(round(elastic_reg_r2_score
elastic_reg_adj_r2_score=1-((1-elastic_reg_r2_score)*(len(y_test)-1)/(len(y_test)-X_test.sh
print("Adjusted R square accuracy is {} %".format(round(elastic_reg_adj_r2_score*100,2)))
```

Our Elatic-Net regression model has 49.812 % accuracy  
Adjusted R square accuracy is 40.07 %

## Comparisions of all Models

### MSE

In [148]:

```
print("MSE for Linear Regression Model is '{}'\nMSE for Ridge Regression Model is '{}'\nMSE
.format(round(mean_squared_error(y_test, linear_reg_pred),2), round(mean_squared_error(y_te
round(mean_squared_error(y_test, lasso_reg_pred),2), round(mean_squared_error(y_test, elas
```

MSE for Linear Regression Model is '5.25'  
MSE for Ridge Regression Model is '5.19'  
MSE for Lasso Regression Model is '6.09'  
MSE for Elastic-Net Regression Model is

In [146]:

```
print("MAE for Linear Regression Model is '{}'\nMAE for Ridge Regression Model is '{}'\nMAE
format(round(mean_absolute_error(y_test, linear_reg_pred),2), round(mean_absolute_error(y_t
round(mean_absolute_error(y_test, lasso_reg_pred),2), round(mean_absolute_error(y_test, elas
```

MAE for Linear Regression Model is '1.81'  
MAE for Ridge Regression Model is '1.8'  
MAE for Lasso Regression Model is '2.0'

